

Time Series Analysis

Reference: <https://www.kaggle.com/competitions/store-sales-time-series-forecasting/data>
(<https://www.kaggle.com/competitions/store-sales-time-series-forecasting/data>)

This is a kaggle competition and in this competition we will use time-series forecasting to forecast store sales on data from Corporación Favorita, a large Ecuadorian-based grocery retailer.

Specifically, we'll build a model that more accurately predicts the unit sales for thousands of items sold at different Favorita stores.

Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [2]: import os
os.listdir('./sample_data/')
```

```
Out[2]: ['holidays_events.csv',
'oil.csv',
'sample_submission.csv',
'stores.csv',
'test.csv',
'train.csv',
'transactions.csv']
```

```
In [3]: holidays_df = pd.read_csv('./sample_data/holidays_events.csv')
oils_df = pd.read_csv('./sample_data/oil.csv')
stores_df = pd.read_csv('./sample_data/stores.csv')
tran_df = pd.read_csv('./sample_data/transactions.csv')
```

Exploring Datasets

holidays_events.csv

Holidays and Events, with metadata

NOTE: Pay special attention to the transferred column.

- A holiday that is transferred officially falls on that calendar day, but was moved to another date by the government.
- A transferred day is more like a normal day than a holiday.

- To find the day that it was actually celebrated, look for the corresponding row where type is Transfer.
- For example, the holiday Independencia de Guayaquil was transferred from 2012-10-09 to 2012-10-12, which means it was celebrated on 2012-10-12.
- Days that are type Bridge are extra days that are added to a holiday (e.g., to extend the break across a long weekend). These are frequently made up by the type Work Day which is a day not normally scheduled for work (e.g., Saturday) that is meant to payback the Bridge.
- Additional holidays are days added a regular calendar holiday, for example, as typically happens around Christmas (making Christmas Eve a holiday).

In [10]: `holidays_df.head()`

Out[10]:

| | date | type | locale | locale_name | description | transferred |
|---|------------|---------|----------|-------------|-------------------------------|-------------|
| 0 | 2012-03-02 | Holiday | Local | Manta | Fundacion de Manta | False |
| 1 | 2012-04-01 | Holiday | Regional | Cotopaxi | Provincializacion de Cotopaxi | False |
| 2 | 2012-04-12 | Holiday | Local | Cuenca | Fundacion de Cuenca | False |
| 3 | 2012-04-14 | Holiday | Local | Libertad | Cantonizacion de Libertad | False |
| 4 | 2012-04-21 | Holiday | Local | Riobamba | Cantonizacion de Riobamba | False |

In [11]: `holidays_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date             350 non-null   object
1   type             350 non-null   object
2   locale           350 non-null   object
3   locale_name      350 non-null   object
4   description       350 non-null   object
5   transferred       350 non-null   bool
dtypes: bool(1), object(5)
memory usage: 14.1+ KB
```

In [12]: `holidays_df.shape`

Out[12]: (350, 6)

```
In [13]: holidays_df.describe().transpose()
```

```
Out[13]:
```

| | count | unique | top | freq |
|-------------|-------|--------|------------|------|
| date | 350 | 312 | 2014-06-25 | 4 |
| type | 350 | 6 | Holiday | 221 |
| locale | 350 | 3 | National | 174 |
| locale_name | 350 | 24 | Ecuador | 174 |
| description | 350 | 103 | Carnaval | 10 |
| transferred | 350 | 2 | False | 338 |

```
In [14]: holidays_df.isna().sum()
```

```
Out[14]: date      0
type      0
locale    0
locale_name 0
description 0
transferred 0
dtype: int64
```

```
In [15]: holidays_df['date'] = pd.to_datetime(holidays_df.date)
holidays_df['month']=holidays_df.date.dt.month
holidays_df['year']=holidays_df.date.dt.year
holidays_df['day']=holidays_df.date.dt.day_name()
holidays_df['day_date']=holidays_df.date.dt.day
holidays_df.head(5)
```

```
Out[15]:
```

| | date | type | locale | locale_name | description | transferred | month | year | day |
|---|------------|---------|----------|-------------|-------------------------------|-------------|-------|------|----------|
| 0 | 2012-03-02 | Holiday | Local | Manta | Fundacion de Manta | False | 3 | 2012 | Friday |
| 1 | 2012-04-01 | Holiday | Regional | Cotopaxi | Provincializacion de Cotopaxi | False | 4 | 2012 | Sunday |
| 2 | 2012-04-12 | Holiday | Local | Cuenca | Fundacion de Cuenca | False | 4 | 2012 | Thursday |
| 3 | 2012-04-14 | Holiday | Local | Libertad | Cantonizacion de Libertad | False | 4 | 2012 | Saturday |
| 4 | 2012-04-21 | Holiday | Local | Riobamba | Cantonizacion de Riobamba | False | 4 | 2012 | Saturday |

oil.csv

- Daily oil price. Includes values during both the train and test data timeframes.
- (Ecuador is an oil-dependent country and it's economical health is highly vulnerable to shocks in oil prices.)

```
In [16]: oils_df.head(5)
```

```
Out[16]:
```

| | date | dcoilwtico |
|---|------------|------------|
| 0 | 2013-01-01 | NaN |
| 1 | 2013-01-02 | 93.14 |
| 2 | 2013-01-03 | 92.97 |
| 3 | 2013-01-04 | 93.12 |
| 4 | 2013-01-07 | 93.20 |

```
In [17]: oils_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1218 entries, 0 to 1217
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        1218 non-null   object
1   dcoilwtico  1175 non-null   float64
dtypes: float64(1), object(1)
memory usage: 19.2+ KB
```

```
In [18]: oils_df.describe().transpose()
```

```
Out[18]:
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------|--------|-----------|-----------|-------|--------|-------|-------|--------|
| dcoilwtico | 1175.0 | 67.714366 | 25.630476 | 26.19 | 46.405 | 53.19 | 95.66 | 110.62 |

```
In [19]: oils_df.isna().sum()
```

```
Out[19]: date        0
dcoilwtico    43
dtype: int64
```

```
In [20]: oils_df.shape
```

```
Out[20]: (1218, 2)
```

```
In [21]: oils_df['date'] = pd.to_datetime(oils_df.date)
oils_df['month']=oils_df.date.dt.month
oils_df['year']=oils_df.date.dt.year
oils_df['day']=oils_df.date.dt.day_name()
oils_df['day_date']=oils_df.date.dt.day
oils_df.head(5)
```

```
Out[21]:
```

| | date | dcoilwtico | month | year | day | day_date |
|---|------------|------------|-------|------|-----------|----------|
| 0 | 2013-01-01 | NaN | 1 | 2013 | Tuesday | 1 |
| 1 | 2013-01-02 | 93.14 | 1 | 2013 | Wednesday | 2 |
| 2 | 2013-01-03 | 92.97 | 1 | 2013 | Thursday | 3 |
| 3 | 2013-01-04 | 93.12 | 1 | 2013 | Friday | 4 |
| 4 | 2013-01-07 | 93.20 | 1 | 2013 | Monday | 7 |

stores.csv

- Store metadata, including city, state, type, and cluster.
- cluster is a grouping of similar stores.

```
In [22]: stores_df.head()
```

```
Out[22]:
```

| | store_nbr | city | state | type | cluster |
|---|-----------|---------------|--------------------------------|------|---------|
| 0 | 1 | Quito | Pichincha | D | 13 |
| 1 | 2 | Quito | Pichincha | D | 13 |
| 2 | 3 | Quito | Pichincha | D | 8 |
| 3 | 4 | Quito | Pichincha | D | 9 |
| 4 | 5 | Santo Domingo | Santo Domingo de los Tsachilas | D | 4 |

```
In [23]: stores_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54 entries, 0 to 53
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   store_nbr   54 non-null    int64
1   city        54 non-null    object
2   state       54 non-null    object
3   type        54 non-null    object
4   cluster     54 non-null    int64
dtypes: int64(2), object(3)
memory usage: 2.2+ KB
```

```
In [24]: stores_df.describe().transpose()
```

```
Out[24]:
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|-----------|-------|-----------|-----------|-----|-------|------|-------|------|
| store_nbr | 54.0 | 27.500000 | 15.732133 | 1.0 | 14.25 | 27.5 | 40.75 | 54.0 |
| cluster | 54.0 | 8.481481 | 4.693395 | 1.0 | 4.00 | 8.5 | 13.00 | 17.0 |

```
In [25]: stores_df.shape
```

```
Out[25]: (54, 5)
```

```
In [26]: stores_df.isna().sum()
```

```
Out[26]: store_nbr    0
city              0
state             0
type              0
cluster           0
dtype: int64
```

transactions.csv

- Gives the information on number of transactions on particular date by store number

```
In [27]: tran_df.head()
```

```
Out[27]:
```

| | date | store_nbr | transactions |
|---|------------|-----------|--------------|
| 0 | 2013-01-01 | 25 | 770 |
| 1 | 2013-01-02 | 1 | 2111 |
| 2 | 2013-01-02 | 2 | 2358 |
| 3 | 2013-01-02 | 3 | 3487 |
| 4 | 2013-01-02 | 4 | 1922 |

```
In [28]: tran_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 83488 entries, 0 to 83487
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   date             83488 non-null object  
1   store_nbr        83488 non-null int64  
2   transactions     83488 non-null int64  
dtypes: int64(2), object(1)
memory usage: 1.9+ MB
```

```
In [29]: ▶ tran_df.describe().transpose()
```

```
Out[29]:
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|--------------|---------|-------------|------------|-----|--------|--------|--------|--------|
| store_nbr | 83488.0 | 26.939237 | 15.608204 | 1.0 | 13.0 | 27.0 | 40.0 | 54.0 |
| transactions | 83488.0 | 1694.602158 | 963.286644 | 5.0 | 1046.0 | 1393.0 | 2079.0 | 8359.0 |

```
In [30]: ▶ tran_df.isna().sum()
```

```
Out[30]: date          0
store_nbr          0
transactions       0
dtype: int64
```

```
In [31]: ▶ tran_df['date'] = pd.to_datetime(tran_df['date'])
tran_df['month']=tran_df.date.dt.month
tran_df['year']=tran_df.date.dt.year
tran_df['day']=tran_df.date.dt.day_name()
tran_df['day_date']=tran_df.date.dt.day
tran_df.head(5)
```

```
Out[31]:
```

| | date | store_nbr | transactions | month | year | day | day_date |
|---|------------|-----------|--------------|-------|------|-----------|----------|
| 0 | 2013-01-01 | 25 | 770 | 1 | 2013 | Tuesday | 1 |
| 1 | 2013-01-02 | 1 | 2111 | 1 | 2013 | Wednesday | 2 |
| 2 | 2013-01-02 | 2 | 2358 | 1 | 2013 | Wednesday | 2 |
| 3 | 2013-01-02 | 3 | 3487 | 1 | 2013 | Wednesday | 2 |
| 4 | 2013-01-02 | 4 | 1922 | 1 | 2013 | Wednesday | 2 |

train.csv

- The training data, comprising time series of features store_nbr, family, and onpromotion as well as the target sales.
- store_nbr identifies the store at which the products are sold.
- family identifies the type of product sold.
- sales gives the total sales for a product family at a particular store at a given date. Fractional values are possible since products can be * sold in fractional units (1.5 kg of cheese, for instance, as opposed to 1 bag of chips).
- onpromotion gives the total number of items in a product family that were being promoted at a store at a given date.


```
In [37]: train_df['date'] = pd.to_datetime(train_df['date'])
train_df['month']=train_df.date.dt.month
train_df['year']=train_df.date.dt.year
train_df['day']=train_df.date.dt.day_name()
train_df['day_date']=train_df.date.dt.day
train_df.head(5)
```

```
Out[37]:
```

| | id | date | store_nbr | family | sales | onpromotion | month | year | day | day_date |
|---|----|------------|-----------|------------|-------|-------------|-------|------|---------|----------|
| 0 | 0 | 2013-01-01 | 1.0 | AUTOMOTIVE | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 1 | 1 | 2013-01-01 | 1.0 | BABY CARE | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 2 | 2 | 2013-01-01 | 1.0 | BEAUTY | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 3 | 3 | 2013-01-01 | 1.0 | BEVERAGES | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 4 | 4 | 2013-01-01 | 1.0 | BOOKS | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |

test.csv

- The test data, having the same features as the training data. You will predict the target sales for the dates in this file.
- The dates in the test data are for the 15 days after the last date in the training data.

```
In [38]: test_df = pd.read_csv('./sample_data/test.csv')
test_df.head(5)
```

```
Out[38]:
```

| | id | date | store_nbr | family | onpromotion |
|---|---------|------------|-----------|------------|-------------|
| 0 | 3000888 | 2017-08-16 | 1 | AUTOMOTIVE | 0 |
| 1 | 3000889 | 2017-08-16 | 1 | BABY CARE | 0 |
| 2 | 3000890 | 2017-08-16 | 1 | BEAUTY | 2 |
| 3 | 3000891 | 2017-08-16 | 1 | BEVERAGES | 20 |
| 4 | 3000892 | 2017-08-16 | 1 | BOOKS | 0 |

In [39]: `test_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28512 entries, 0 to 28511
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   id              28512 non-null  int64
1   date            28512 non-null  object
2   store_nbr       28512 non-null  int64
3   family          28512 non-null  object
4   onpromotion     28512 non-null  int64
dtypes: int64(3), object(2)
memory usage: 1.1+ MB
```


In [40]: `test_df.shape`

Out[40]: (28512, 5)

In [41]: `test_df.describe().transpose()`

Out[41]:

| | count | mean | std | min | 25% | 50% | 7 |
|-------------|---------|--------------|-------------|-----------|------------|-----------|---------|
| id | 28512.0 | 3.015144e+06 | 8230.849774 | 3000888.0 | 3008015.75 | 3015143.5 | 3022271 |
| store_nbr | 28512.0 | 2.750000e+01 | 15.586057 | 1.0 | 14.00 | 27.5 | 41 |
| onpromotion | 28512.0 | 6.965383e+00 | 20.683952 | 0.0 | 0.00 | 0.0 | 6 |



In [42]: `test_df.isna().sum()`

Out[42]:

| | |
|-------------|---|
| id | 0 |
| date | 0 |
| store_nbr | 0 |
| family | 0 |
| onpromotion | 0 |

dtype: int64

```
In [43]: test_df['date'] = pd.to_datetime(test_df['date'])
test_df['month']=test_df.date.dt.month
test_df['year']=test_df.date.dt.year
test_df['day']=test_df.date.dt.day_name()
test_df['day_date']=test_df.date.dt.day
test_df.head(5)
```

```
Out[43]:
```

| | id | date | store_nbr | family | onpromotion | month | year | day | day_date |
|---|---------|------------|-----------|------------|-------------|-------|------|-----------|----------|
| 0 | 3000888 | 2017-08-16 | 1 | AUTOMOTIVE | 0 | 8 | 2017 | Wednesday | 1 |
| 1 | 3000889 | 2017-08-16 | 1 | BABY CARE | 0 | 8 | 2017 | Wednesday | 1 |
| 2 | 3000890 | 2017-08-16 | 1 | BEAUTY | 2 | 8 | 2017 | Wednesday | 1 |
| 3 | 3000891 | 2017-08-16 | 1 | BEVERAGES | 20 | 8 | 2017 | Wednesday | 1 |
| 4 | 3000892 | 2017-08-16 | 1 | BOOKS | 0 | 8 | 2017 | Wednesday | 1 |

Exploratory Data Analysis

- Holidays


In [44]: `holidays_df.head(10)`

Out[44]:


| | date | type | locale | locale_name | description | transferred | month | year | day |
|---|------------|---------|----------|-------------|-------------------------------|-------------|-------|------|----------|
| 0 | 2012-03-02 | Holiday | Local | Manta | Fundacion de Manta | False | 3 | 2012 | Friday |
| 1 | 2012-04-01 | Holiday | Regional | Cotopaxi | Provincializacion de Cotopaxi | False | 4 | 2012 | Sunday |
| 2 | 2012-04-12 | Holiday | Local | Cuenca | Fundacion de Cuenca | False | 4 | 2012 | Thursday |
| 3 | 2012-04-14 | Holiday | Local | Libertad | Cantonizacion de Libertad | False | 4 | 2012 | Saturday |
| 4 | 2012-04-21 | Holiday | Local | Riobamba | Cantonizacion de Riobamba | False | 4 | 2012 | Saturday |
| 5 | 2012-05-12 | Holiday | Local | Puyo | Cantonizacion del Puyo | False | 5 | 2012 | Saturday |
| 6 | 2012-06-23 | Holiday | Local | Guaranda | Cantonizacion de Guaranda | False | 6 | 2012 | Saturday |
| 7 | 2012-06-25 | Holiday | Regional | Imbabura | Provincializacion de Imbabura | False | 6 | 2012 | Monday |
| 8 | 2012-06-25 | Holiday | Local | Latacunga | Cantonizacion de Latacunga | False | 6 | 2012 | Monday |
| 9 | 2012-06-25 | Holiday | Local | Machala | Fundacion de Machala | False | 6 | 2012 | Monday |




```
In [45]: ▶ type_order = holidays_df['type'].value_counts().index
sns.countplot(data=holidays_df, x='type', order=type_order)
plt.title('Holidays based on types')
plt.show()
```




```
In [46]: ▶ sns.countplot(data=holidays_df, x='locale')  
plt.title('Holidays based on location types')  
plt.show()
```




```
In [47]: ▶ l_order = holidays_df['locale_name'].value_counts().index
plt.figure(figsize=(15,10))
sns.countplot(data=holidays_df, x='locale_name',order=l_order)
plt.title('Holidays based on locations')
plt.xticks(rotation=45)
plt.show()
```



```
In [48]: ▶ sns.countplot(data=holidays_df[holidays_df.type=='Holiday'],x='day')  
plt.title("Holiday based on day")  
plt.show()
```




```
In [49]: ▶ sns.countplot(data=holidays_df[holidays_df.type=='Holiday'],x='day', hue='t
plt.title("Transferred Holiday based on day")
plt.show()
```



- Oils

```
In [50]: ▶ oils_df.head(10)
```

```
Out[50]:
```


| | date | dcoilwtico | month | year | day | day_date |
|---|------------|------------|-------|------|-----------|----------|
| 0 | 2013-01-01 | NaN | 1 | 2013 | Tuesday | 1 |
| 1 | 2013-01-02 | 93.14 | 1 | 2013 | Wednesday | 2 |
| 2 | 2013-01-03 | 92.97 | 1 | 2013 | Thursday | 3 |
| 3 | 2013-01-04 | 93.12 | 1 | 2013 | Friday | 4 |
| 4 | 2013-01-07 | 93.20 | 1 | 2013 | Monday | 7 |
| 5 | 2013-01-08 | 93.21 | 1 | 2013 | Tuesday | 8 |
| 6 | 2013-01-09 | 93.08 | 1 | 2013 | Wednesday | 9 |
| 7 | 2013-01-10 | 93.81 | 1 | 2013 | Thursday | 10 |
| 8 | 2013-01-11 | 93.60 | 1 | 2013 | Friday | 11 |
| 9 | 2013-01-14 | 94.27 | 1 | 2013 | Monday | 14 |

```
In [51]: oils_mean = oils_df.dcoilwtico.mean()
oils_df.dcoilwtico.fillna(oils_mean,inplace=True)
oils_df.isna().sum()
```

```
Out[51]: date          0
dcoilwtico          0
month              0
year              0
day              0
day_date          0
dtype: int64
```

```
In [52]: sns.distplot(oils_df['dcoilwtico'], kde=True, bins=25)
```

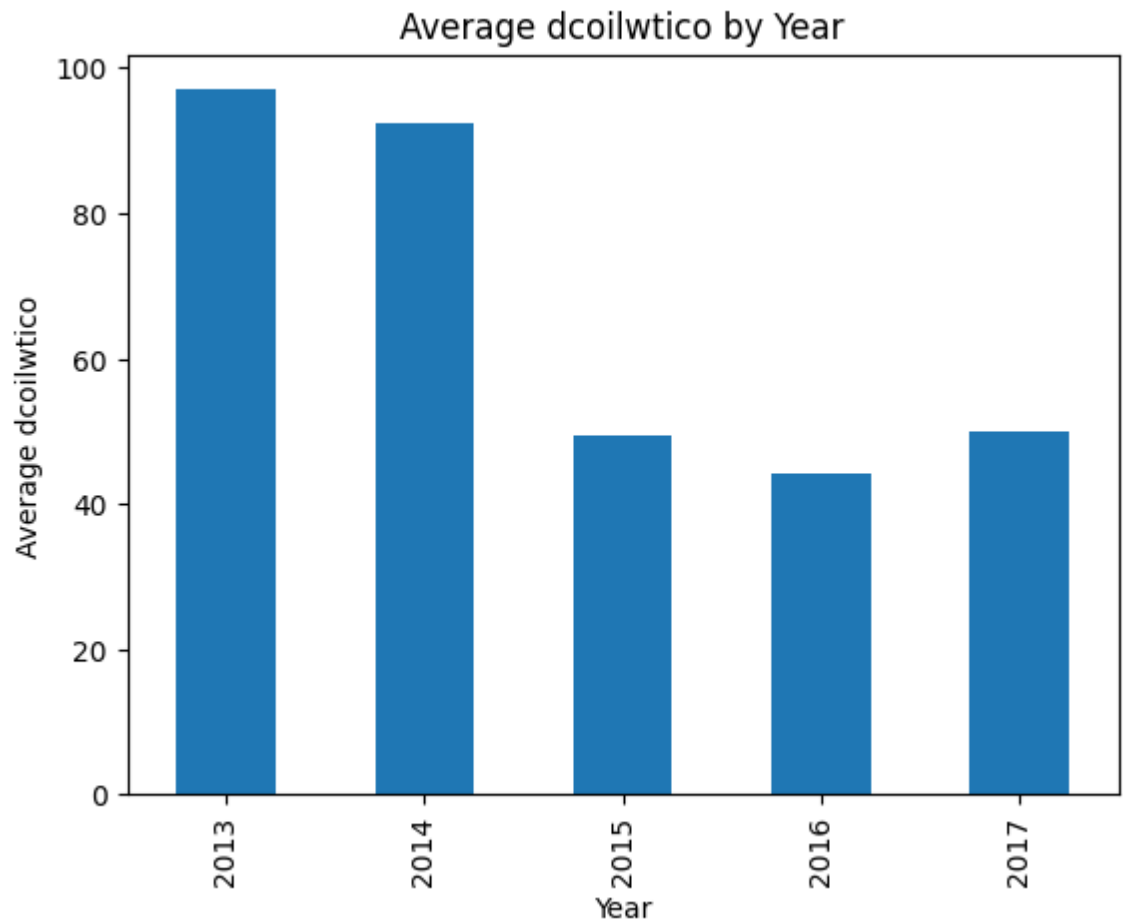
```
Out[52]: <Axes: xlabel='dcoilwtico', ylabel='Density'>
```




```
In [53]: oils_df.groupby('year')['dcoilwtico'].mean()
```

```
Out[53]: year
2013    96.938810
2014    92.294365
2015    49.313867
2016    44.135744
2017    49.976383
Name: dcoilwtico, dtype: float64
```

```
In [54]: pivot_table = oils_df.pivot_table(index='year', values='dcoilwtico', aggfun  
  
# Plotting the graph  
pivot_table.plot(kind='bar', legend=False)  
plt.xlabel('Year')  
plt.ylabel('Average dcoilwtico')  
plt.title('Average dcoilwtico by Year')  
plt.show()
```



```
In [55]: oils_year_month = oils_df.pivot_table(index='month', columns='year', values
# You can separate data with lines
sns.heatmap(oils_year_month, cmap='Blues', linecolor='white', linewidth=1)
plt.title('Oil produced by year for every month')
plt.show()
```




```
In [56]: stores_df.head(10)
```

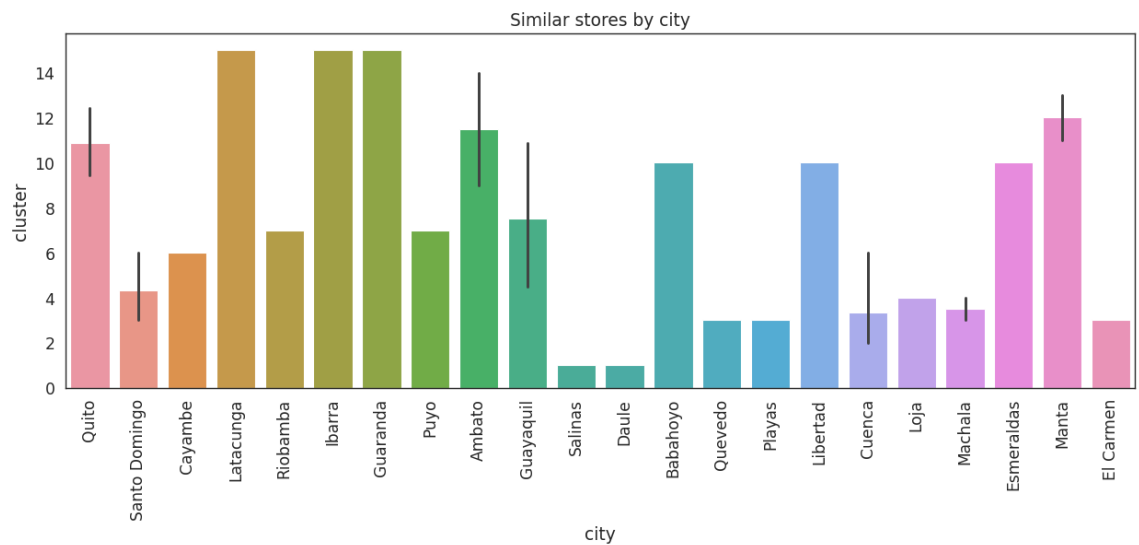
```
Out[56]:
```

| | store_nbr | city | state | type | cluster |
|---|-----------|---------------|--------------------------------|------|---------|
| 0 | 1 | Quito | Pichincha | D | 13 |
| 1 | 2 | Quito | Pichincha | D | 13 |
| 2 | 3 | Quito | Pichincha | D | 8 |
| 3 | 4 | Quito | Pichincha | D | 9 |
| 4 | 5 | Santo Domingo | Santo Domingo de los Tsachilas | D | 4 |
| 5 | 6 | Quito | Pichincha | D | 13 |
| 6 | 7 | Quito | Pichincha | D | 8 |
| 7 | 8 | Quito | Pichincha | D | 8 |
| 8 | 9 | Quito | Pichincha | B | 6 |
| 9 | 10 | Quito | Pichincha | C | 15 |


```
In [96]: ▶ plt.figure(figsize=(15,5))
sns.barplot(x='state',y='cluster',data=stores_df)
plt.xticks(rotation=90)
plt.title("Similar stores by state")
plt.show()
```



```
In [95]: ▶ plt.figure(figsize=(15,5))
sns.barplot(x='city',y='cluster',data=stores_df)
plt.xticks(rotation=90)
plt.title("Similar stores by city")
plt.show()
```



```
In [59]: ▶ s_order = stores_df['type'].value_counts().index
sns.countplot(x='type',data=stores_df,order=s_order)
plt.title('Number of stores by type')
plt.show()
```




```
In [60]: ▶ stores_sales_df=pd.merge(train_df,stores_df,on='store_nbr',how='inner')
stores_sales_df.head(5)
```


Out[60]:

| | id | date | store_nbr | family | sales | onpromotion | month | year | day | day_date |
|---|----|------------|-----------|------------|-------|-------------|-------|------|---------|----------|
| 0 | 0 | 2013-01-01 | 1.0 | AUTOMOTIVE | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 1 | 1 | 2013-01-01 | 1.0 | BABY CARE | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 2 | 2 | 2013-01-01 | 1.0 | BEAUTY | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 3 | 3 | 2013-01-01 | 1.0 | BEVERAGES | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |
| 4 | 4 | 2013-01-01 | 1.0 | BOOKS | 0.0 | 0.0 | 1 | 2013 | Tuesday | 1 |

```
In [61]: ▶ sales_by_sate=stores_sales_df.groupby('state').sales.sum().to_frame()
plt.figure(figsize=(15,5))
sns.barplot(data=sales_by_sate,x=sales_by_sate.sales.index,y=sales_by_sate.
plt.title('Sales by State')
plt.xticks(rotation=90)
plt.show()
```



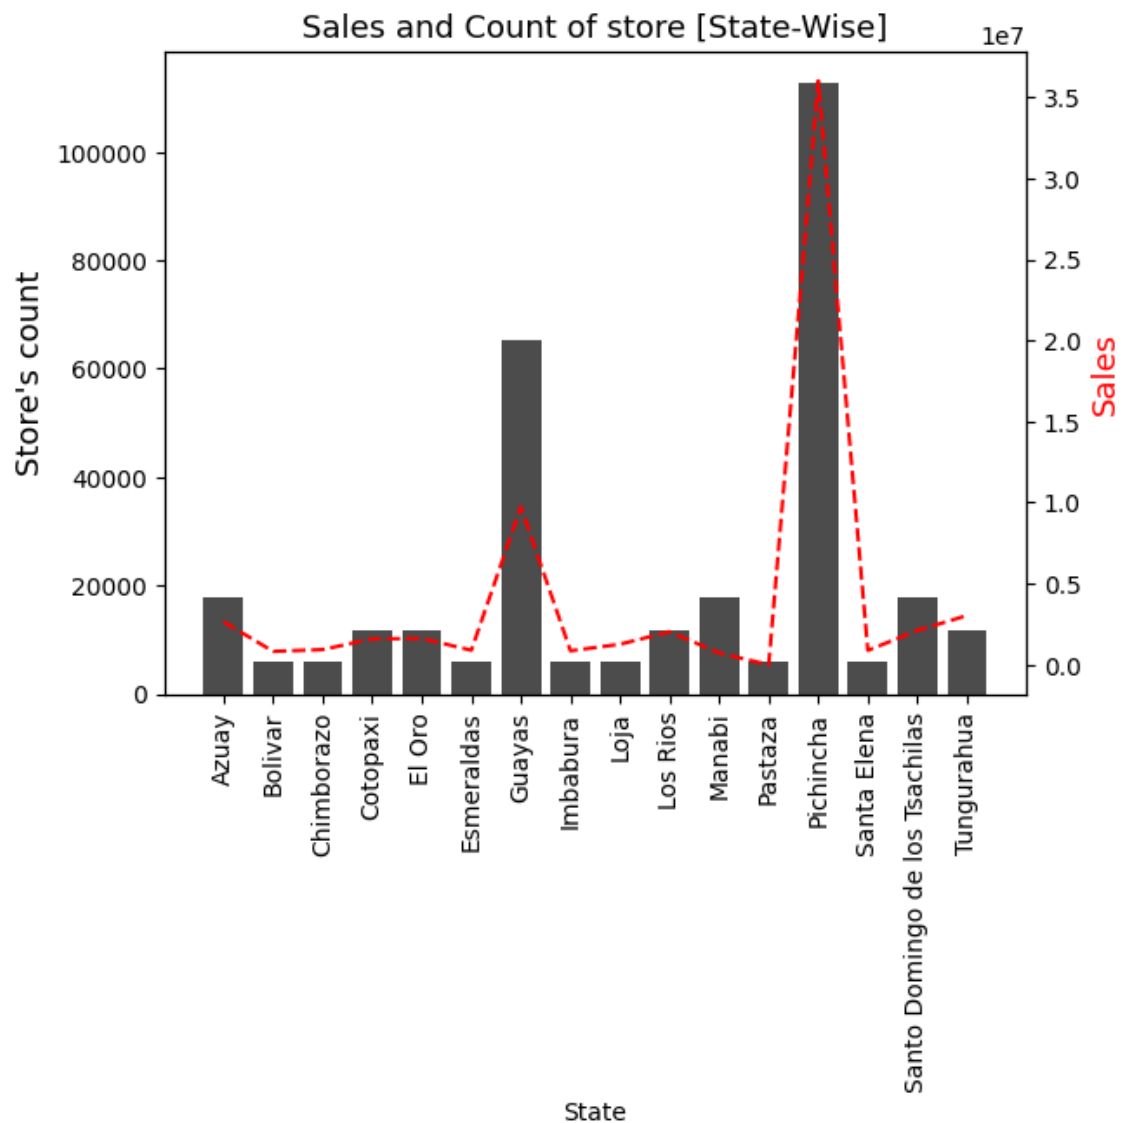
```
In [62]: ▶ sales_by_city=stores_sales_df.groupby('city').sales.sum().to_frame()
plt.figure(figsize=(15,5))
sns.barplot(data=sales_by_city,x=sales_by_city.sales.index,y=sales_by_city.
plt.title('Sales by City')
plt.xticks(rotation=90)
plt.show()
```



```
In [63]: ▶ # Creating list based of requirements i.e (name and value should be on same
state=stores_sales_df.groupby('state')
count_store=state.store_nbr.count()
states=[s for s,df in state]
sales_state=state.sales.sum()

# Ploting
fig,ax1=plt.subplots()
ax2=ax1.twinx()
ax1.bar(states,count_store,color='black',alpha=0.7)
ax2.plot(states,sales_state,'r--')

ax1.set_xlabel("State")
ax1.set_ylabel("Store's count",color='black',size=13)
ax2.set_ylabel("Sales",color='red',size=13)
ax1.set_xticklabels(states,rotation='vertical',size=10)
plt.title("Sales and Count of store [State-Wise]",size=13)
plt.show()
```




- Transactions

```
In [64]: ▶ tran_df.head()
```


```
Out[64]:
```

| | date | store_nbr | transactions | month | year | day | day_date |
|---|------------|-----------|--------------|-------|------|-----------|----------|
| 0 | 2013-01-01 | 25 | 770 | 1 | 2013 | Tuesday | 1 |
| 1 | 2013-01-02 | 1 | 2111 | 1 | 2013 | Wednesday | 2 |
| 2 | 2013-01-02 | 2 | 2358 | 1 | 2013 | Wednesday | 2 |
| 3 | 2013-01-02 | 3 | 3487 | 1 | 2013 | Wednesday | 2 |
| 4 | 2013-01-02 | 4 | 1922 | 1 | 2013 | Wednesday | 2 |

```
In [65]: ▶ sns.histplot(data=tran_df,x='transactions',kde=True,bins=12)  
plt.title("Transactions Count")  
plt.show()
```

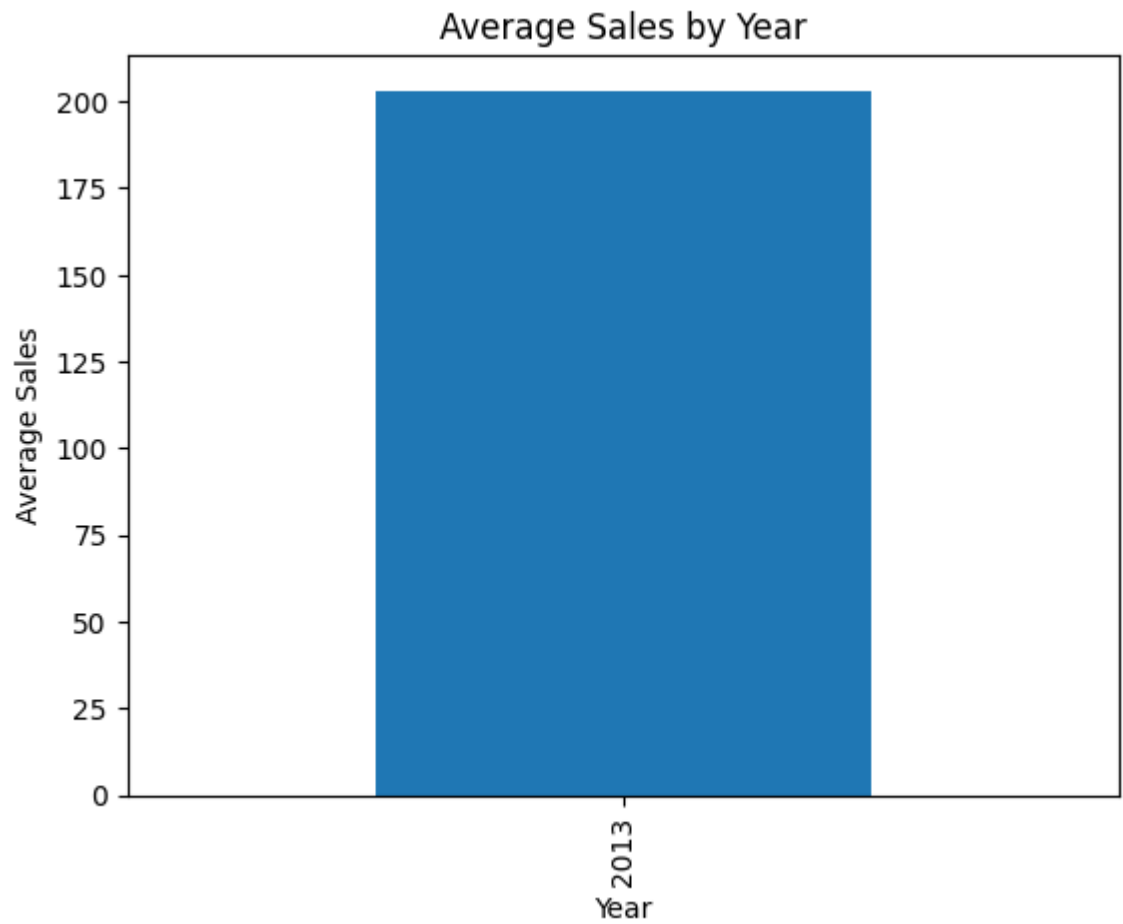


```
In [66]: ▶ transactions_year_month = tran_df.pivot_table(index='month', columns='year'  
# You can separate data with lines  
sns.heatmap(transactions_year_month, cmap='Blues', linecolor='white', linewidth=1,  
plt.title('Transactions done by year for every month')  
plt.show()
```




- Train DataFrame

```
In [67]: ▶ pivot_table = train_df.pivot_table(index='year', values='sales', aggfunc='  
  
# Plotting the graph  
pivot_table.plot(kind='bar', legend=False)  
plt.xlabel('Year')  
plt.ylabel('Average Sales')  
plt.title('Average Sales by Year')  
plt.show()
```




```
In [68]: ▶ dowpivot_table = train_df.pivot_table(index='day', values='sales', aggfunc=  
# Plotting the graph  
dowpivot_table.plot(kind='bar', legend=False)  
plt.xlabel('Day of the week')  
plt.ylabel('Average Sales')  
plt.title('Average Sales by Day of Week')  
plt.show()
```



```
In [69]: pf_pivot_table = train_df.pivot_table(index='family', values='sales', aggfu

# Plotting the graph
plt.figure(figsize=(15,5))
pf_pivot_table.plot(kind='bar', legend=False)
plt.xlabel('Product Category')
plt.ylabel('Average Sales')
plt.title('Average Sales by Product Category')
plt.show()
```

<Figure size 1500x500 with 0 Axes>



Feature Engineering

```
In [72]: ▶ print("Train Data set columns : ",list(train_df.columns))
print("Test Data set columns : ",list(test_df.columns))
print("Store Data set columns : ",list(stores_df.columns))

Train Data set columns : ['id', 'date', 'store_nbr', 'family', 'sales',
'onpromotion', 'month', 'year', 'day', 'day_date']
Test Data set columns : ['id', 'date', 'store_nbr', 'family', 'onpromotio
n', 'month', 'year', 'day', 'day_date']
Store Data set columns : ['store_nbr', 'city', 'state', 'type', 'cluste
r']
```

```
In [73]: ▶ trainDf = pd.merge(train_df,stores_df,how='inner',on='store_nbr')
testDf = pd.merge(test_df,stores_df,how='inner',on='store_nbr')
print("Train Data set columns : ",list(trainDf.columns))
print("Test Data set columns : ",list(testDf.columns))

Train Data set columns : ['id', 'date', 'store_nbr', 'family', 'sales',
'onpromotion', 'month', 'year', 'day', 'day_date', 'city', 'state', 'typ
e', 'cluster']
Test Data set columns : ['id', 'date', 'store_nbr', 'family', 'onpromotio
n', 'month', 'year', 'day', 'day_date', 'city', 'state', 'type', 'cluste
r']
```

Defining the input and target columns

```
In [74]: ▶ input_cols = ['store_nbr', 'family', 'onpromotion', 'month', 'year', 'day',
target_cols = ['sales']]
```

```
In [75]: ▶ #obtaining the non object dtypes
numerical_features = [feature for feature in trainDf[input_cols] if trainDf
print(numerical_features)
categorical_features = [feature for feature in trainDf[input_cols] if train
print(categorical_features)

['store_nbr', 'onpromotion', 'month', 'year', 'day_date', 'cluster']
['family', 'day', 'city', 'state', 'type']
```

```
In [76]: ▶ print(trainDf[numerical_features].isna().sum())
print(trainDf[categorical_features].isna().sum())
```

```
store_nbr      0
onpromotion    0
month          0
year           0
day_date       0
cluster        0
dtype: int64
family         0
day            0
city           0
state          0
type           0
dtype: int64
```

Data Preprocessing

```
In [77]: ▶ trainDf[numerical_features]
```

Out[77]:

| | store_nbr | onpromotion | month | year | day_date | cluster |
|--------|-----------|-------------|-------|------|----------|---------|
| 0 | 1.0 | 0.0 | 1 | 2013 | 1 | 13 |
| 1 | 1.0 | 0.0 | 1 | 2013 | 1 | 13 |
| 2 | 1.0 | 0.0 | 1 | 2013 | 1 | 13 |
| 3 | 1.0 | 0.0 | 1 | 2013 | 1 | 13 |
| 4 | 1.0 | 0.0 | 1 | 2013 | 1 | 13 |
| ... | ... | ... | ... | ... | ... | ... |
| 320762 | 9.0 | 0.0 | 6 | 2013 | 29 | 6 |
| 320763 | 9.0 | 0.0 | 6 | 2013 | 29 | 6 |
| 320764 | 9.0 | 0.0 | 6 | 2013 | 29 | 6 |
| 320765 | 9.0 | 0.0 | 6 | 2013 | 29 | 6 |
| 320766 | 9.0 | 0.0 | 6 | 2013 | 29 | 6 |

320767 rows × 6 columns

```
In [78]: ▶ from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(trainDf[numerical_features])
```

Out[78]: StandardScaler()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [79]: trainDf[numerical_features]=scaler.transform(trainDf[numerical_features])
         testDf[numerical_features]=scaler.transform(testDf[numerical_features])
```

Encoding for categorical features

```
In [80]: from sklearn.preprocessing import OneHotEncoder
```

```
In [81]: ▶ encoder=OneHotEncoder(sparse_output=False,handle_unknown='ignore')
encoder.fit(trainDf[categorical_features])
```

```
Out[81]: OneHotEncoder(handle_unknown='ignore', sparse_output=False)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [82]: ► encoded_cols=list(encoder.get_feature_names_out(categorical_features))
```

```
In [83]: trainDf[encoded_cols]=encoder.transform(trainDf[categorical_features])
          testDf[encoded_cols]=encoder.transform(testDf[categorical_features])
```

```
In [84]: ▶ trainData=trainDf[numerical_features+encoded_cols]
testData=testDf[numerical_features+encoded_cols]
trainData.shape,testData.shape
```

```
Out[84]: ((320767, 89), (28512, 89))
```

Model Building

```
In [85]: from sklearn.model_selection import train_test_split
```

```
In [86]: X = trainData.copy()
          y = trainDf[target_cols]
```

```
In [87]: X_train, X_test, y_train, y_test = train_test_split(X,y ,
random_state=104,
test_size=0.25,
shuffle=True)
```



```

In [94]: ▶ from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

# Models to evaluate
models = [
    ('Linear Regression', LinearRegression()),
    ('Decision Tree Regressor', DecisionTreeRegressor(random_state=42)),
    # ('Random Forest Regressor', RandomForestRegressor(random_state=42))
]

# Evaluate models using k-fold cross-validation
for model_name, model in models:
    # For RandomForestRegressor, add hyperparameters to the parameter grid
    if model_name == 'Random Forest Regressor':
        param_grid = {
            'n_estimators': [10, 15, 30],
            'max_depth': [5, 10, 20],
            'min_samples_split': [2, 5, 8],
            'min_samples_leaf': [1, 2, 4]
        }
    else:
        # For other models, use a basic parameter grid
        param_grid = {}

    # Create the GridSearchCV object
    grid_search = GridSearchCV(model, param_grid, cv=5, scoring='neg_mean_squared_error')

    # Fit the model to the training data
    grid_search.fit(X_train, y_train)

    # Get the best parameters and the best model
    best_params = grid_search.best_params_
    best_model = grid_search.best_estimator_

    # Print the best parameters
    print(f"\nBest Parameters for {model_name}: ", best_params)

    # Evaluate the best model using k-fold cross-validation
    mse_scores = cross_val_score(best_model, X_train, y_train, cv=5, scoring='neg_mean_squared_error')
    rmse_scores = np.sqrt(-mse_scores)

    # Print cross-validation results
    print(f"Cross-validation RMSE scores for {model_name}: {rmse_scores}")
    print(f"Mean RMSE: {np.mean(rmse_scores)}")

    # Make predictions on the test set
    y_pred = best_model.predict(X_test)

    # Evaluate the model on the test set
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    print(f"\nTest set RMSE for {model_name}: {rmse}")

    # Optionally, you can print other metrics as well, like R2 score

```

```
from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print(f"R2 score for {model_name}: {r2}")
```

Best Parameters for Linear Regression: {}
Cross-validation RMSE scores for Linear Regression: [448.43794289 434.2791
5129 431.68155695 427.60886844 446.64514099]
Mean RMSE: 437.73053211292955

Test set RMSE for Linear Regression: 437.92352023957966
R2 score for Linear Regression: 0.5888385989674367

Best Parameters for Decision Tree Regressor: {}
Cross-validation RMSE scores for Decision Tree Regressor: [179.7441305 18
1.23371422 214.30101129 173.26191091 221.32767169]
Mean RMSE: 193.9736877194814

Test set RMSE for Decision Tree Regressor: 182.30486476295638
R2 score for Decision Tree Regressor: 0.9287454977477645

In [88]: ▶